

Exploiting Tunability of Erasure Coding for Low-Interference Repair

Paper ID: 27

1 Summary of Implementations

We implement a prototype of ChameleonEC from scratch with around 3,000 lines of code (LoC) in C++. We employ Jerasure Library v2.0 coupled with GF-complete to enable the encoding and decoding functionalities over Galois Field arithmetic. We also employ the interfaces of Redis to realize the data transmission and storage across nodes. We elaborate on the dependent software and testing tools as follows.

Dependent software: The implementation of ChameleonEC relies on a suite of software, including Jerasure Library v2.0, GF-complete, Redis, and nethogs, which can be reached via the following URLs.

```
#!/bin/bash
$ wget https://github.com/intel/intel-
-cmt-cat
$ git clone https://github.com/ceph/
gf-complete.git
$ git clone https://github.com/
tsuraan/Jerasure.git
$ wget https://github.com/raboof/
nethogs/archive/v0.8.5.tar.gz
```

Testing tools: we employ YCSB, IBM Object Store, Memcached(Twitter) (a trace released by Twitter), and Facebook ETC workloads in our evaluation. They can be reached via the following URLs.

```
#!/bin/bash
$ wget https://github.com/
brianfrankcooper/YCSB/releases/
download/0.17.0/ycsb-0.17.0.tar.gz
$ http://iotta.snia.org/traces/
key-value/28652
$ http://iotta.snia.org/traces
/key-value/36305
$ https://dl.acm.org/doi/pdf/
10.1145/2254756.2254766
```

```
<settings>
<attribute><name>erasure.code.type</name><value>RS</value></attribute>
<attribute><name>erasure.code.k</name><value>18</value></attribute>
<attribute><name>erasure.code.n</name><value>14</value></attribute>
<attribute><name>node.fail.cn</name><value>1</value></attribute>
<attribute><name>jrc.code.1</name><value>0</value></attribute>
<attribute><name>period.len</name><value>20</value></attribute>
<attribute><name>encode.matrix.file</name><value>/root/ElasticRepair/conf/rsencMat_10_14</value></attribute>
<attribute><name>packet.size</name><value>1048576</value></attribute>
<attribute><name>packet.count</name><value>64</value></attribute>
<attribute><name>single.stripe.alg</name><value>path</value></attribute>
<attribute><name>schedule.method</name><value>online</value></attribute>
<attribute><name>coordinator.address</name><value>172.31.0.201</value></attribute>
<attribute><name>file.system.type</name><value>standalone</value></attribute>
<attribute><name>meta.stripe.dir</name><value>/root/repairboost-code/meta/standalone-meta</value></attribute>
<attribute><name>file.system.type</name><value>standalone</value></attribute>
<attribute><name>block.directory</name><value>/root/repairboost-code/meta/standalone-blocks</value></attribute>
<attribute><name>helpers.address</name>
<value>default/172.31.0.202</value>
<value>default/172.31.0.203</value>
<value>default/172.31.0.204</value>
<value>default/172.31.0.205</value>
<value>default/172.31.0.206</value>
<value>default/172.31.0.207</value>
<value>default/172.31.0.208</value>
<value>default/172.31.0.209</value>
<value>default/172.31.0.210</value>
<value>default/172.31.0.211</value>
<value>default/172.31.0.212</value>
<value>default/172.31.0.213</value>
<value>default/172.31.0.214</value>
<value>default/172.31.0.215</value>
<value>default/172.31.0.216</value>
<value>default/172.31.0.217</value>
<value>default/172.31.0.218</value>
<value>default/172.31.0.219</value>
<value>default/172.31.0.220</value>
</attribute>
<attribute><name>local.ip.address</name><value>172.31.0.201</value></attribute>
</settings>
```

Figure 1: Configurations of ChameleonEC.

2 Evaluation Details

Hardware configurations: We assess the performance of ChameleonEC on Amazon EC2. We allocate 20 virtual machine instances with the type of m5.xlarge in the region of North Virginia. Each instance runs Ubuntu 16.04.7 LTS and owns four vCPU with 3.1 GHz Intel Xeon Platinum processor, 16 GB RAM, and 100 GB Elastic Block Store (EBS) volumes. The network bandwidth between any two instances is around 5 Gb/s (measured via iperf).

We employ the tool wondershaper to throttle the network bandwidth, which is set to 1 Gb/s by default.

```
# !/bin/bash
$ wondershaper -a ens5 -d 1048576 -u 1048576
```

We perform extensive testbed experiments, including (i) the experiments to understand the properties of ChameleonEC; and (ii) the experiments to understand the sensitivity of ChameleonEC.

How to start ChameleonEC: We run the prototype of ChameleonEC on Amazon EC2 to assess its performance.

```

<property>
  <name>hbase.cluster.distributed</name>
  <value>true</value>
</property>
<property>
  <name>hbase.rootdir</name>
  <value>hdfs://Master:9000/hbase</value>
</property>
<property>
  <name>hbase.zookeeper.quorum</name>
  <value>Master,Slave2,Slave3,Slave4,Slave5,Slave6,Slave7,Slave8,Slave9,Slave10,
  Slave11,Slave12,Slave13,Slave14,Slave15,Slave16,Slave17,Slave18,Slave19,Slave20</value>
</property>
<property>
  <name>hbase.zookeeper.property.dataDir</name>
  <value>/root/zookeeper</value>
</property>
</configuration>

```

Figure 2: Major configurations of Hbase.

Figure 1 shows the major configurations of our testbed. We use the following commands to start ChameleonEC:

We use the script `run-BD-Monitor.py` to start the bandwidth monitor.

```
#!/bin/bash
$ python3 run-BD-Monitor.py
```

We run the `ECCoordinator` on the Master to start the ChameleonEC coordinator.

```
#!/bin/bash
$ ./ECCoordinator
```

We run the `ECHelper` on the slaves to start the ChameleonEC helpers.

```
#!/bin/bash
$ ./ECHelper
```

We run the `ECClient` on the failed node to send a repair request.

```
#!/bin/bash
$ ./ECClient
```

We write a script `start.py` to start `ECCoordinator` and `ECHelper` conveniently and use `stop.py` to kill them.

```
#!/bin/bash
$ python3 start.py
$ python3 stop.py
```

Experiments with YCSB: We measure the impact of network bandwidth, chunk sizes, slice sizes, and erasure coding parameters under YCSB. YCSB relies on Hadoop and Hbase. We omit their installation process here. Figure 2 shows the major configurations of Hbase.

The following script starts Hadoop and Hbase to run YCSB:

```
#!/bin/bash
$ start-all.sh
$ start-hbase.sh
```

We use YCSB WorkloadA by default. Figure 3 shows the configurations of YCSB WorkloadA.

We use the following command to load data and run the YCSB server:

```
recordcount=200000
operationcount=200000
workload=site.ycsb.workloads.CoreWorkload
```

```
readallfields=true
```

```
readproportion=0.50
updateproportion=0.50
scanproportion=0
insertproportion=0
```

```
requestdistribution=zipfian
```

```
table=usertable
columnfamily=family
```

```
fieldlength=65536
fieldcount=1
```

Figure 3: Configurations of YCSB WorkloadA.

```
#!/bin/bash
$ ./root/ycsb-0.17.0/bin/ycsb load hbase20
-P /root/ycsb-0.17.0 /workloads/testworkload
-cp /root/hbase-2.3.6/conf/ -threads 2
```

```
$ ./root/ycsb-0.17.0/bin/ycsb run hbase20
-P /root/ycsb-0.17.0 /workloads/testworkload
-cp /root/hbase-2.3.6/conf/ -threads 2
```

Experiments with Memcached: We also employ IBM Object Store, Memcached (Twitter), and Facebook ETC workloads to clarify how we assess the performance of ChameleonEC. The experiments with Memcached workloads are similar. We use IBM Object Store as an example. For every instance, we run the following commands to start Memcached:

```
#!/bin/bash
$ service Memcached stop
$ memcached -d install -m 8192 -u root -p
11211 -c 4096 -I 1M
```

We use `IBM-load` and `IBM-run` to load data and start IBM Object Store server.

```
#!/bin/bash
$ ./IBM-load
$ ./IBM-run
```

We write a script `IBM-run.py` to perform these operations. We also use `Facebook-run.py`, `Twitter-run.py` to run other Memcached servers.

```
#!/bin/bash
$ python3 IBM-run.py 0 (run IBM trace only)
$ python3 IBM-run.py 1 (run ChameleonEC only)
$ python3 IBM-run.py 2 (run IBM trace and ChameleonEC)
$ python3 IBM-run.py 3 (initialize and load IBM trace
data)
```